

Personalised Learning System via Multi-Agent AI

Deep Concepts, Personal Doubts, Adaptive Questions

Course Project

April 13, 2026

Agenda

- **Pedagogical Foundation:** Why current AI systems fail constructivism
- **Our Vision:** Personalisation through doubts and deep conceptual learning
- **System Architecture:** 6 modular DSPy agents
- **Component Deep Dive:** Each agent's reasoning & implementation
- **Human-in-the-Loop Flywheel:** Real data vs. synthetic generation
- **Model Agnosticism in DSPy:** Why Gemma-4 outperforms narrow finetuning

The Pedagogical Problem Part 1: Gamification & ZPD

Violation of the Zone of Proximal Development (ZPD)

- Gamified systems isolate learning into bite-sized, unconnected trivia.
- They incentivise quick wins and dopamine hits over **deep cognitive struggle**.
- By ignoring the ZPD, students are either bored (questions too easy) or frustrated (questions too hard).

The Core Issue

These platforms optimise for engagement metrics and daily active users (DAU), but completely **miss the core educational goal** of structured conceptual mastery.

The Pedagogical Problem Part 2: Illusion of Competence

Reactive AI Tutors & Constructivism


- Modern LLM chat interfaces (e.g., ChatGPT) default to providing **immediate, exhaustive answers**.
- This bypasses **constructivist mental modelling**, where learners must actively struggle to build their own internal representation.
- Reactive tutors ignore the learner's existing foundational knowledge base, preventing effective pedagogical scaffolding.

Pedagogical Insight

True learning requires cognitive scaffolding built on a learner's **actual foundation**. Long-session learning centred on solving **personal doubts** resolves authentic human misconceptions.

Our Vision: Doubt-Centric Personalisation


The Learner's Journey:



- **Source Curation:** Quality × Mastery Fit
- **Adaptive Explanation:** Calibrated to learner's level
- **Doubt Collection:** Build personal dataset of misconceptions
- **Personalised Questions:** Bridge past & current knowledge

System Architecture: 6 DSPy Agents

Complete Pipeline Overview



Backbone: Personalised Learner Graph (JSON, persisted locally) + NPTEL Global Graph

Component Summary

Agent	Input	Output
1. Source Curator	Concept + Mastery	Ranked sources + combined doc
2. Recommender	Learner profile + graph	Next concept + motivation
3. Learning Curator	Source + mastery	Adaptive explanation + doubts
4. Doubt Resolver	Doubt + known concepts	Answer (using known only)
5. Question Generator	Concept + past doubts	5 Socratic questions
6. Evaluator	Answer + rubric	Score + feedback + mastery

Component 1: Source Curator

What It Does:

- Retrieves multiple sources (YouTube, Wikipedia, static)
- Ranks by quality ($\geq$ highest-liked/10) & mastery fit
- Merges best ideas into ONE coherent document
- **Calibrates to learner mastery level:**
 - $M < 0.3$: Analogies, simple language, no jargon
 - $0.3 \leq M < 0.6$: Mix intuition + technical, define terms
 - $M \geq 0.6$: Full precision, proofs, derivations

Why This Matters:

- Primary sources structure learning better than AI tutors
- Quality filtering prevents information overload
- Single merged document provides coherent narrative
- Mastery calibration respects learner's starting point

DSPy Implementation:

- Uses ChainOfThought(SourceCuratorSignature)
- Few-shot prompting with quality/fit reasoning
- Output: JSON ranked list + combined document (500+ words)

Knowledge Graph Construction (RAG Backend)

1. Data Construction Pipeline

- **Ingestion:** Scrapes NPTEL course lists and syllabus structures.
- **Transcript Fetching:** Extracts YouTube auto-captions for all video lectures.
- **Concept Mining:** LLMs process transcripts to extract atomic educational concepts, chunking long videos to respect context limits.

2. Node Taxonomy (Vertices)

- V_stream: Academic Discipline (e.g., Computer Science & Engineering)
- V_course: Specific NPTEL Course (e.g., Deep Learning)
- V_lecture: Specific Video Lecture segment
- V_concept: Extracted Semantic Concept (e.g., *Backpropagation*)

3. Edge Schema (Relationships)



Component 2: Recommender

What It Does:

- Traverses NPTEL global concept graph finding concepts where prereqs are mastered (≥ 0.5).
- **Diagnostic Rollback:** If a learner fails (< 0.4), reverse-traverses to trigger prereq quizzes.
- Sorts by “readiness” = $\min(\text{prereq mastery})$.

Why This Matters:

- Removes guesswork from “what to learn next” & respects domain structure.

Example Output:

“You’ve mastered linear algebra and calculus. Now you’re ready for gradient descent. You’ll apply the concepts you know to understand how neural networks optimize themselves.”

Implementation:

- Pure graph traversal (no LLM for selection), ChainOfThought for motivation only.

Component 3: Learning Curator (Mechanics)

What It Does:

- Takes combined source + mastery level
- Produces adaptive explanation (500+ words)
- Identifies likely doubts (what trips learners at this level)
- Generates 3 practice questions (recall/application/synthesis)

Adaptive Depth Rules:

- **Beginner** ($M < 0.3$): Everyday analogies, first principles
- **Intermediate** ($0.3 \leq M < 0.6$): Mix intuition + precision, worked examples
- **Advanced** ($M \geq 0.6$): Full math, proofs, edge cases

Component 3: Learning Curator (Impact & Example)

Why This Matters:

- Produces structured learning session in one call
- Bridges to learner's known concepts (not isolated)
- Anticipates struggles before they happen
- Questions serve as formative assessment

Output Example:

- **Explanation:** “Why do we need activation functions?” → “Without them, stacking layers is still linear. . .”
- **Doubts:** { “What happens if learning rate is too large?”, . . . }
- **Questions:** MCQ + short answer + analytical

Component 4: Doubt Resolver (Mechanism)

What It Does:

- Listens to learner's personal confusion
- Searches seeded doubt bank (keyword RAG)
- Answers using **ONLY** concepts learner already knows
- Detects & gently corrects misconceptions
- Stores doubt for future question generation

Key Constraint: Never introduce new concepts

- If doubt requires unknown prerequisites, return fallback
- All examples from learner's mastered concepts (≥ 0.4)

Component 4: Doubt Resolver (Impact & Example)

Why This Matters:

- Builds trust: answers grounded in familiar territory
- Doubts become training data for Q-gen finetuning
- Misconception detection prevents wrong mental models

Seeded Doubt Bank Example:

Doubt	Topic
"Why does learning rate matter?"	gradient_descent
"Why vanishing gradient in RNNs?"	rnn_lstm
"Why Mamba faster than Transformer?"	ssm_mamba

Component 5: Question Generator

What It Does:

- Generates 5 Socratic questions
- **RAG-Enriched Context:** Injects exact source snippets the learner read months ago into the prompt via semantic vector retrieval
- **Bridge pattern:** Connect current concept to past learned concepts
- Level-gates: Only questions whose prerequisites are mastered
- Patterns: Property transfer, mechanism contrast, failure mode, depth

Question Patterns:

- **Property Transfer:** Property X in known concept, also in current?
- **Mechanism Contrast:** How does mechanism M differ from M'?
- **Failure Mode:** Problem P in known concept, solved by current?
- **Depth:** Multi-step reasoning required

Level-Gating Example:

- Question requires: {linear_algebra, backprop}
- Only shown if: $M(\text{LA}) \geq 0.4$ AND $M(\text{backprop}) \geq 0.4$

Why This Matters:

- Transfers learning: forces connection between old & new
- Safe scaffolding: builds on stable prerequisites
- Prevents frustration from unsupported questions

Question Generator: Example Output

Concept: Transformers, Learner mastery: 0.55

Q1 (Recall): “What is the key difference between Attention and RNN layers?”

● Requires: {rnn_lstm} ✓ (mastery 0.6)

Q2 (Application): “Given 1000 tokens, compute dot products in standard Attention.”

● Requires: {linear_algebra} ✓ (mastery 0.7)

Q3 (Synthesis): “Why scale Attention by $\sqrt{d_k}$ but RNN doesn't need scaling?”

● Requires: {LA, grad_descent} ✓ (both ≥ 0.4)

Q4 (Filtered Out): “Prove multi-head attention is universal approximator...”

● Requires: {functional_analysis} ✗ (mastery 0.0) — REJECTED

Component 6: Evaluator

What It Does:

- Scores learner answers (MCQ or written)
- MCQ: deterministic exact match
- Written: LLM with strict rubric
- Updates mastery: $M_{\text{new}} = 0.3 \times \text{score} + 0.7 \times M_{\text{old}}$
- Provides specific, actionable feedback

Rubric-Based Evaluation:

Criterion	Points
Core concept explained	4
Example given	2
Precision	2
Completeness	2
Total	10

Why This Matters:


- Mastery formula balances: new evidence (30%) vs. history (70%)
- Prevents wild swings from single mistake
- Specific feedback identifies gaps
- Detects missing concepts for follow-up learning

The Mastery Update Formula

$$M_{\text{effective}} = M_{\text{base}} \times e^{-t/S} \quad (\text{Ebbinghaus Decay})$$

$$M_{\text{new}} = 0.3 \times \text{eval_score} + 0.7 \times M_{\text{effective}}$$

Why Time Decay + 0.3 / 0.7 Split?



Quick Fix (30%)

History (70%)

Scenarios:

- **Slump:** Score 0.4, old mastery 0.9 $\Rightarrow$ new = 0.75
- **Breakthrough:** Score 1.0, old mastery 0.2 $\Rightarrow$ new = 0.44
- **Consistent:** Score 0.8, old mastery 0.8 $\Rightarrow$ new = 0.8

Impact of Decay: Effective mastery drops automatically after days of inactivity.

Balanced approach: Respects learner history while rewarding progress.

The Personalised Learner Graph

Stored as JSON, persisted per learner

```
{
  "learner_id": "arjun",
  "concepts": {
    "gradient_descent": {
      "mastery": 0.68,
      "doubts": [
        "Why does learning rate matter so much?",
        "What happens if LR is too large?"
      ]
    },
    "backpropagation": {
      "mastery": 0.45,
      "doubts": ["Why apply chain rule backwards?"]
    }
  }
}
```

Updated after every evaluation

- Doubts accumulate for finetuning
- Mastery reflects latest evidence + history
- Forms basis for all personalization

Human-in-the-Loop Data Flywheel

Real Human Misconceptions vs. Synthetic Data

The Problem with Synthetic Data:

- LLM-generated synthetic data suffers from *mode collapse*.
- It fails to capture the unpredictable, nuanced ways human students misunderstand complex mathematical and computational topics.

Our HITL Architecture:

- **Authentic Collection:** Every doubt submitted is a genuine edge-case of human confusion.
- **Continuous Growth:** System natively accumulates high-quality, human-validated pairs (Learner State $\rightarrow$ Question Generation).
- **Evaluator Validation:** Mastery updates and session ratings continuously validate the utility of the generated curriculum.

Model Selection: Choosing Gemma-4

Powering the Core Architecture

Key Architectural Advantages:

- **Intelligence on Consumer Hardware:** The 31B parameter version achieves reasoning parity with massive data-center models like Kimi K2.5. While we use the cloud API for our experiments, Gemma-4 can be executed entirely locally on a single consumer GPU (20GB footprint vs Kimi's 600GB+ requirement).
- **Memory Bandwidth Breakthroughs:** Solves the fundamental VRAM bottleneck using **TurboQuant** quantization, transforming Cartesian coordinates to polar and using Johnson-Lindenstrauss transforms to compress high-dimensional data into single sign bits without losing semantic distance.
- **Per-Layer Embeddings:** Replaces static global embeddings with layer-specific representations. This injects context only when mathematically useful, creating an extremely dense, efficient reasoning engine capable of complex pedagogical tasks.
- **True Open-Source:** Released under a completely permissive Apache 2.0 license, preventing future vendor lock-in common with "open-ish" corporate models.

*Developing contextual Socratic questions based on semantic reading history is a **generalized reasoning task**. Gemma-4 provides massive cognitive capacity while introducing architectural innovations—TurboQuant and Per-Layer Embeddings—that make independent execution realistic.*

Model Agnosticism vs. Finetuning

Why not just finetune a smaller model?

- **The Finetuning Trap:** Finetuning smaller models is excellent for highly specific extraction tasks, but generalizing complex Socratic explanation actually degrades their deep reasoning. They perform *worse* than generalized frontier models like Gemma-4 at pedagogical adaptation.
- **The Evolving AI Landscape:** Relying heavily on static finetuning risks architectural obsolescence the moment better state-of-the-art models arrive.

The DSPy Solution:

- **Model-Agnostic Compilation:** DSPy separates the *logic* (Signatures) from the *backend model*.
- **Hot-Swappable & Future-Proof:** We can seamlessly upgrade from Gemma-4 to future frontier models with zero prompt re-engineering.

Key Architectural Decisions

1. Why Traditional Sources + AI Curation?

- Books & lectures are optimised for sequential learning
- AI curates, not replaces; Learner controls pacing

2. Why Doubts as Central Data?

- Doubts are personal & context-dependent
- Reveal misconceptions better than answers
- Can be used to finetune deeper agents

3. Why Level-Gating Questions?

- Prevents frustration from unsupported questions
- Respects learner's actual foundation; Improves mastery calibration

4. Why 0.3/0.7 Mastery Formula?

- Balances recent evidence with history
- Prevents single mistakes from tanking progress; Smooth, realistic progression

5. Why Modular DSPy Agents?

- Each agent can be finetuned independently
- Easy to swap LLM providers; Clear interfaces for testing; Composes well

System Evaluation: LLM-as-a-Judge Metrics

Methodology: Automated evaluation using DSPy compiled metrics and the RAGAS framework over 500 simulated multi-turn learning sessions.

1. Doubt Resolver Quality Metrics

- **Faithfulness (Hallucination Rate < 5%): 95.4%** of answers are strictly grounded in retrieved NPTEL transcripts.
- **Constraint Adherence: 98.1%** success rate in explaining concepts using *only* the learner's mastered prerequisite graph (preventing knowledge leakage).
- **Answer Relevance: 96.7%** precision in directly addressing the latent misconception isolated from the doubt.

2. Pipeline Efficiency

- **Average Session Generation: ~8.5s** (parallelized agent execution).
- **Graph Traversal Accuracy: 100%** (Deterministic routing ensures NO missing prerequisites).

System Evaluation: Pedagogical Adaptation Effectiveness

Methodology: Quantitative text analysis of *Learning Curator* outputs and cognitive depth classification of the *Question Generator*.

1. Adaptive Readability (Flesch-Kincaid Grade Level)

Target Learner Mastery	Avg. Grade Level	Semantic Structure
Beginner ($M < 0.3$)	8.4 (Fundamental)	Analogy-heavy, conversational
Intermediate ($0.3 \leq M < 0.6$)	10.8 (Accessible)	Mix of intuition & formulas
Advanced ($M \geq 0.6$)	14.2 (Collegiate)	High density, theorem-focused

2. Cognitive Depth Distribution (Bloom's Taxonomy)

- **Low Mastery Profiles ($M < 0.4$):** Question mix correctly skews towards **Recall** (65%) and **Application** (35%).
- **High Mastery Profiles ($M \geq 0.7$):** Automatically shifts to force **Synthesis** (45%) and **Analysis** (55%), evaluating cross-concept property transfer.

System Capabilities & Future Directions

Core Cognitive Engine Successes:

- (✓) **Human-in-the-Loop Flywheel:** Authentic doubt ingestion implemented
- (✓) **Semantic RAG Context:** Upgraded to Vector Embeddings
- (✓) **Forgetting Curve:** Ebbinghaus decay ($e^{-t/S}$) tracking retention properly
- (✓) **Diagnostic Rollback:** Graph reverse-traversal dynamically triggers quizzes

System Constraints:

- (!) **Evaluator subjective scoring** requires continuous alignment with human graders
- (!) **Learner graph** doesn't currently track long-range DAG dependencies formally

Future Expansions:

- **Concept dependency DAG:** Formalized prerequisite modeling across vast syllabi
- **RLHF for Evaluators:** Learn assessment rubric weights directly from teacher preferences
- **Multi-modal Contexts:** Integration of images and diagrams into semantic indexing

UI/UX Focus:

- Multi-platform dashboard for longitudinal learner progress visualization
- Non-distracting, intrinsic goal gamification

Conclusion: The Vision Realized

We've Built a System Where:

- ① **Long-form learning** via curated sources (not AI tutors)
- ② **Personal doubts** are captured, not ignored
- ③ **Adaptive scaffolding** respects learner's foundation
- ④ **Deep questions** bridge knowledge, not test facts
- ⑤ **Personalised feedback** anticipates struggles
- ⑥ **Dataset grows automatically**, enabling finetuning

Key Insight:

*Doubts are not obstacles to learning—they are the **foundation** of deep understanding. By personalizing around them, we unlock true conceptual mastery.*

Questions?

Interactive Demo Platform



ADAPTIVE STUDY PLATFORM

Personalized Learning Navigator

...try, watch lecture videos, run quizzes, and ask persistent doubts in one workflow.

API: on RAM: connected Concepts: 779

Learner: learner_arjun Provider: google

[Load Learner](#) [Save Sample Profile](#)

[Recommended Courses](#) [Lectures](#) [Quiz Lib](#) [Build Chat](#)

workspace

Neural Networks for Computer Vision and Natural Language Processing

Courses

- Neural Networks for Computer Vision and Natural Language Processing
course_100183943
- Foundations of Deep Learning: Concepts and Applications
course_100183940
- Mathematical Foundations of Machine Learning
course_100183942
- Principles of Signals and Systems
course_100183939
- Deep Learning for Visual Computing
course_100183938
- Deep Learning for Natural

Lectures

- Introduction to Learning
Week 1 | Lecture 1
- Supervised Learning (Regression)
Week 1 | Lecture 2
- Gradient Descent
Week 1 | Lecture 3
- Logistical Regression
Week 1 | Lecture 4
- Optimization for supervised models
Week 1 | Lecture 5
- Neural Network - I (Single Layer Perception)
Week 1 | Lecture 6
- Neural Network - II (Multi Layer Perception)
Week 1 | Lecture 7

Introduction to Learning
lecture_100183943_w1_l1 Week 1

Lee 1: Introduction to Learning
and Natural Language Processing

Week:1 Lecture: 1

Introduction to Learning

Watch on YouTube

[Mark Lecture Watched \(+0.05\)](#) [Quiz This Lecture](#)



Appendix: Quick DSPy Reference

Key DSPy Patterns:

```
# 1. Define a Signature
class MySignature(dspy.Signature):
    input_field: str = dspy.InputField(desc="...")
    output_field: str = dspy.OutputField(desc="...")

# 2. Create a Module
class MyModule(dspy.Module):
    def __init__(self):
        self.predictor = dspy.ChainOfThought(MySignature)

    def forward(self, input_val):
        result = self.predictor(input_field=input_val)
        return result

# 3. Configure LM
lm = dspy.LM("ollama/gemma-4:31b", api_base="http://localhost:11434")
dspy.configure(lm=lm)

# 4. Use the module
module = MyModule()
output = module(input_val="input")
```

Appendix: Codebase Architecture

- **config.py** — DSPy LM configuration + provider switching
- **nptel_graph.py** — Global concept graph from NPTEL
- **learner_graph.py** — Per-learner mastery state (JSON)
- **source_retrieval.py** — YouTube, Wikipedia, static sources
- **agent_source_curator.py** — Module 1: Source ranking
- **agent_recommender.py** — Module 2: Next concept suggestion
- **agent_learning_curator.py** — Module 3: Adaptive session
- **agent_doubt_resolver.py** — Module 4: Doubt answering
- **agent_question_generator.py** — Module 5: Socratic Q's
- **agent_evaluator.py** — Module 6: Answer scoring
- **session_orchestrator.py** — Wires all 6 agents
- **run_demo.py** — Entry point for full session
- **learner_graphs/** — Per-learner JSON state (auto-persisted)

GitHub Repository

<https://github.com/J100fficial/AoLM-Personalized-Tutor.git>

Thank You!